

Predictive geometry claims are sensitive to the comparator

Post-result sensitivity: all 14 existing features, four recipe baselines and two displacement-conditioned comparators.

	B0 RMSE 2.905 p = 8	B1 RMSE 0.652 p = 11	B2 RMSE 1.391 p = 21	B3 RMSE 1.439 p = 24	B0 + disp. RMSE 1.004 p = 9	B3 + disp. RMSE 1.420 p = 25
Log segment displacement / weight	-0.59*	+0.04	-0.03*	-0.05*	+0.24	-0.04*
Segment stable rank	+0.11	-0.13*	+0.04	+0.16	+0.05	+0.14
Segment entropy rank (top-64)	-0.52*	+0.02	-0.01*	-0.00*	+0.22	-0.00*
Segment row-norm CV	-0.53*	-0.02*	+0.01	+0.01	+0.14	+0.01
Segment top-1% row energy	-0.49*	-0.09*	+0.01	+0.01	+0.03	+0.01
Cumulative displacement / weight	-1.90*	+0.14	+0.02	-0.02*	-0.00	-0.00
Cumulative stable rank	+0.35	+0.42	+0.05	+0.12	+0.24	+0.14
Segment overlap with AdamW	+0.24	-0.00	+0.15	+0.24	-0.19*	+0.23
Cumulative overlap with AdamW	+0.84	+0.37	+0.23	+0.49	+0.26	-0.08
Exact segment stable rank	+0.11	-0.12*	+0.05	+0.19	+0.05	+0.16
Segment entropy rank (full spectrum)	-0.56*	+0.04	-0.10*	-0.08*	+0.26	-0.07*
Exact cumulative stable rank	+0.35	+0.35	-0.08*	-0.07*	+0.21	-0.04
Exact segment overlap with AdamW	+0.16	-0.05	+0.12	+0.21	-0.16*	+0.18
Exact cumulative overlap with AdamW	+0.41	+0.75	+0.07	+0.70	+0.45	+0.60

Negative values favor the added feature. *: lower pooled RMSE and improvements in $\geq 3/4$ held-dose folds.
B0: original; B1: + raw LR; B2: optimizer $\times$ stage + log/raw LR; B3: B2 + nominal cumulative LR budget.
Same 60 states and training seed. All flags are exploratory, not significance or causal evidence.